

Going Further with the parallel quicksort

Benjamin Cathelineau

26/11/2021

Modifying the shell script in order to acquire metadata about the condition when the experiment was started

As suggested, I added a lot of metadata, like the number of processes, the cpu and ram usage...

I looked on the internet to find a way to extract most of these metadata.

```
cat scripts/run_benchmarking2.sh
```

```
##
## OUTPUT_DIRECTORY=data/'hostname'_'date +%F'
## mkdir -p $OUTPUT_DIRECTORY
## OUTPUT_FILE=$OUTPUT_DIRECTORY/measurements_'date +%R'.csv
##
## touch $OUTPUT_FILE
##
## # scheduler settings
## sudo renice --priority -19 --pid $BASHPID > /dev/null
## echo "Size, TimeSeq, TimePar,TimeBuiltIn, CPU_USAGE, RAM_USAGE, STORAGE_USAGE, N_PROCESS, ON_BATTERY, TEMPERATURE" >> $OUTPUT_FILE
## for i in $(seq 1 5); do
##     for j in $(seq 0 1000000 10000000) ; do
##         echo -n "$j" >> $OUTPUT_FILE
##         ./src/parallelQuicksort $j >> $OUTPUT_FILE
##     done
##
##     # Collection metadata
##     # cpu usage https://unix.stackexchange.com/questions/69167/bash-script-that-print-cpu-usage-etc
##     echo -n ',top -b -n1 | grep "Cpu(s)" | awk '{print $2 + $4}'' >> $OUTPUT_FILE
##
##     # ram usage
##     FREE_DATA='free -m | grep Mem'
##     CURRENT='echo $FREE_DATA | cut -f3 -d' ' '
##     TOTAL='echo $FREE_DATA | cut -f2 -d' ' '
##     echo -n ',$(echo "scale = 2; $CURRENT/$TOTAL*100" | bc)>> $OUTPUT_FILE
##
##     # hdd usage
##     echo -n ',df -lh | awk '{if ($6 == "/") { print $5 }}' | head -1 | cut -d'%' -f1'>> $OUTPUT_FILE
##
##     # number of processes
##     echo -n ',$(ps aux | wc -l)>> $OUTPUT_FILE
##
##     # on battery or not
##     echo -n ', >> $OUTPUT_FILE
##     upower -i 'upower -e | grep 'BAT'' | grep 'state' | sed "s/state://g" | tr -d '\n' | xargs | tr -d ' ' >> $OUTPUT_FILE
##     # Temperature https://askubuntu.com/questions/779819/cpu-temperature-embedded-in-bash-command-prompt
##     echo -n ',$(sensors | grep -oP 'CPU Die Core Temp.*?\+\K[0-9.]+')>> $OUTPUT_FILE
```

```
##
## # cpu governor https://unix.stackexchange.com/questions/182696/how-to-get-current-cpupower-governor
## echo ,$(cat /sys/devices/system/cpu/cpu0/cpufreq/scaling_governor)>> $OUTPUT_FILE
##     done
## done
```

Modyfing the C program in order not to have to deal with perl

Now, the C program directly outputs a csv file, because I don't know Perl and I don't have time to learn it. The printing was changed.

```
printf(", %lf ", diff);
```

Customizing the linux scheduler settings in order to increase our test priorities

At the start of the shell script, I run this command:

```
sudo renice --priority -19 --pid $BASHPID > /dev/null
```

This changes the scheduler setting so that our test is favored. See the man page. In short a lower niceness is better for the process. A higher niceness is worse.

Analysing the new data, with confidence intervals

```
library(dplyr)
```

```
##
## Attachement du package : 'dplyr'
## Les objets suivants sont masqués depuis 'package:stats':
##
##     filter, lag
## Les objets suivants sont masqués depuis 'package:base':
##
##     intersect, setdiff, setequal, union
```

```
library(ggplot2)
library(Rmisc)
```

```
## Le chargement a nécessité le package : lattice
## Le chargement a nécessité le package : plyr
## -----
## You have loaded plyr after dplyr - this is likely to cause problems.
## If you need functions from both plyr and dplyr, please load plyr first, then dplyr:
## library(plyr); library(dplyr)
## -----
##
## Attachement du package : 'plyr'
## Les objets suivants sont masqués depuis 'package:dplyr':
##
##     arrange, count, desc, failwith, id, mutate, rename, summarise,
##     summarize
```

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --

## v tibble 3.1.6      v purrr 0.3.4
## v tidyr 1.1.4      v stringr 1.4.0
## v readr 2.1.1      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x plyr::arrange() masks dplyr::arrange()
## x purrr::compact() masks plyr::compact()
## x plyr::count() masks dplyr::count()
## x plyr::failwith() masks dplyr::failwith()
## x dplyr::filter() masks stats::filter()
## x plyr::id() masks dplyr::id()
## x dplyr::lag() masks stats::lag()
## x plyr::mutate() masks dplyr::mutate()
## x plyr::rename() masks dplyr::rename()
## x plyr::summarise() masks dplyr::summarise()
## x plyr::summarize() masks dplyr::summarize()

df = read.csv("data/fedora_2021-12-01/measurements_14:37.csv") # load the data from the CSV
df
```

	Size	TimeSeq	TimePar	TimeBuiltIn	CPU_USAGE	RAM_USAGE	STORAGE_USAGE
## 1	0	0.000000	0.000511	0.000004	9	46	14
## 2	1000000	0.122787	0.180089	0.201578	10	46	14
## 3	2000000	0.258838	0.345271	0.394852	10	46	14
## 4	3000000	0.389309	0.472851	0.614174	9	46	14
## 5	4000000	0.537799	0.643878	0.830870	25	46	14
## 6	5000000	0.713201	0.801830	1.074060	17	46	14
## 7	6000000	0.818768	0.902943	1.297154	11	46	14
## 8	7000000	1.006788	1.042709	1.533210	11	46	14
## 9	8000000	1.181180	1.255406	1.767254	11	46	14
## 10	9000000	1.257162	1.410159	2.031089	12	46	14
## 11	10000000	1.451663	1.634807	2.303282	12	46	14
## 12	0	0.000001	0.000371	0.000003	14	46	14
## 13	1000000	0.125794	0.193168	0.190120	13	46	14
## 14	2000000	0.262169	0.335554	0.421010	19	46	14
## 15	3000000	0.409404	0.505379	0.662281	10	46	14
## 16	4000000	0.567011	0.675450	0.903268	15	46	14
## 17	5000000	0.700046	0.830256	1.126786	18	46	14
## 18	6000000	0.869244	0.938737	1.361411	10	46	14
## 19	7000000	0.985444	1.176316	1.599060	14	46	14
## 20	8000000	1.132755	1.327023	1.788308	17	46	14
## 21	9000000	1.309040	1.492488	2.211324	16	46	14
## 22	10000000	1.430431	1.605811	2.503851	11	46	14
## 23	0	0.000000	0.000505	0.000002	12	46	14
## 24	1000000	0.126942	0.195896	0.190241	13	46	14
## 25	2000000	0.268907	0.338938	0.407934	22	46	14
## 26	3000000	0.423347	0.510146	0.669175	15	46	14
## 27	4000000	0.591148	0.660241	0.886063	16	46	14
## 28	5000000	0.715351	0.807237	1.098474	36	46	14
## 29	6000000	0.838489	0.957679	1.343032	23	46	14
## 30	7000000	0.989882	1.116152	1.673185	10	46	14
## 31	8000000	1.120642	1.242341	1.817161	20	46	14

## 32	9000000	1.315064	1.401234	1.994647	10	46	14
## 33	10000000	1.407706	1.575461	2.216965	9	46	14
## 34	0	0.000000	0.000372	0.000011	15	46	14
## 35	1000000	0.127388	0.183451	0.197383	21	46	14
## 36	2000000	0.272571	0.335810	0.397416	8	46	14
## 37	3000000	0.391157	0.481363	0.613656	13	46	14
## 38	4000000	0.539626	0.648876	0.837274	11	46	14
## 39	5000000	0.670290	0.789260	1.055956	12	46	14
## 40	6000000	0.811055	0.957940	1.300443	10	46	14
## 41	7000000	0.953955	1.094826	1.538587	11	46	14
## 42	8000000	1.123812	1.281357	1.763244	18	46	14
## 43	9000000	1.276537	1.431894	1.965513	9	46	14
## 44	10000000	1.400998	1.536765	2.238027	21	46	14
## 45	0	0.000000	0.000539	0.000003	18	46	14
## 46	1000000	0.123637	0.185461	0.197273	11	46	14
## 47	2000000	0.256073	0.331379	0.413583	9	46	14
## 48	3000000	0.402250	0.512991	0.608351	9	46	14
## 49	4000000	0.536034	0.651125	0.839474	10	46	14
## 50	5000000	0.671789	0.781475	1.055982	11	46	14
## 51	6000000	0.850489	0.942595	1.321626	14	46	14
## 52	7000000	0.978512	1.096399	1.533914	27	46	14
## 53	8000000	1.152610	1.368908	2.070410	22	46	14
## 54	9000000	1.308921	1.407782	2.117979	11	46	14
## 55	10000000	1.499842	1.915462	2.464082	10	46	14
##	N_PROCESS	ON_BATTERY	TEMPERATURE	CPU_GOVERNOR			
## 1	374	fully-charged	62.2	schedutil			
## 2	374	fully-charged	63.0	schedutil			
## 3	375	fully-charged	67.0	schedutil			
## 4	375	fully-charged	72.2	schedutil			
## 5	375	fully-charged	70.5	schedutil			
## 6	375	fully-charged	74.2	schedutil			
## 7	375	fully-charged	76.0	schedutil			
## 8	375	fully-charged	77.0	schedutil			
## 9	376	fully-charged	78.5	schedutil			
## 10	375	fully-charged	81.2	schedutil			
## 11	375	fully-charged	82.0	schedutil			
## 12	375	fully-charged	82.0	schedutil			
## 13	375	fully-charged	83.2	schedutil			
## 14	375	fully-charged	84.8	schedutil			
## 15	375	fully-charged	86.2	schedutil			
## 16	375	fully-charged	83.8	schedutil			
## 17	375	fully-charged	86.8	schedutil			
## 18	375	fully-charged	87.5	schedutil			
## 19	375	fully-charged	89.0	schedutil			
## 20	376	fully-charged	91.8	schedutil			
## 21	376	fully-charged	86.8	schedutil			
## 22	376	fully-charged	87.2	schedutil			
## 23	376	fully-charged	84.8	schedutil			
## 24	376	fully-charged	84.8	schedutil			
## 25	376	fully-charged	88.8	schedutil			
## 26	376	fully-charged	88.0	schedutil			
## 27	376	fully-charged	90.5	schedutil			
## 28	376	fully-charged	87.8	schedutil			
## 29	376	fully-charged	87.8	schedutil			

```
## 30      376 fully-charged      87.2      schedutil
## 31      376 fully-charged      89.5      schedutil
## 32      376 fully-charged      89.5      schedutil
## 33      376 fully-charged      85.5      schedutil
## 34      376 fully-charged      85.5      schedutil
## 35      376 fully-charged      87.0      schedutil
## 36      376 fully-charged      89.2      schedutil
## 37      376 fully-charged      88.5      schedutil
## 38      376 fully-charged      91.5      schedutil
## 39      376 fully-charged      91.2      schedutil
## 40      376 fully-charged      89.2      schedutil
## 41      376 fully-charged      89.0      schedutil
## 42      376 fully-charged      89.8      schedutil
## 43      377 fully-charged      88.5      schedutil
## 44      378 fully-charged      88.2      schedutil
## 45      378 fully-charged      88.2      schedutil
## 46      377 fully-charged      87.0      schedutil
## 47      377 fully-charged      88.8      schedutil
## 48      377 fully-charged      90.5      schedutil
## 49      377 fully-charged      88.0      schedutil
## 50      377 fully-charged      86.5      schedutil
## 51      378 fully-charged      88.5      schedutil
## 52      377 fully-charged      90.5      schedutil
## 53      377 fully-charged      89.0      schedutil
## 54      378 fully-charged      92.5      schedutil
## 55      380 fully-charged      87.5      schedutil
```

You can see that now there is a lot more metadata for each entry

plot and confidence interval on the time


Here I plot the new data and get essentially the same result as before. I have to do a wierd transformation because of the way the data is now.

```
# https://www.datanovia.com/en/fr/blog/comment-creer-un-ggplot-contenant-plusieurs-lignes/
df2 <- df %>%
  dplyr::select(Size, TimeSeq, TimePar, TimeBuiltIn) %>%
  gather(key = "variable", value = "value", -Size)

p <- ggplot(df2, aes(x=Size, y = value, color = variable)) + geom_point() +geom_smooth()

print(p)

## 'geom_smooth()' using method = 'loess' and formula 'y ~ x'
```



Then I try to compute the confidence interval on time for each size and for each algorithm. That's a lot of data.

```
# Intended for a .95 confidence
MY_CI <- function(data){
  data_mean = mean(data)
  std_dev = sd(data)
  final_mult= std_dev/length(data)
  return(c(up= data_mean+(final_mult*2),mean=data_mean,down= data_mean-(final_mult*2)))
}
```

```
df %>% group_by(Size) %>% group_map(~ CI(x=.x$TimeSeq,ci=.95))
```

```
## [[1]]
##      upper      mean      lower
## 7.55289e-07 2.00000e-07 -3.55289e-07
##
## [[2]]
##      upper      mean      lower
## 0.1278223 0.1253096 0.1227969
##
## [[3]]
##      upper      mean      lower
## 0.2722683 0.2637116 0.2551549
##
## [[4]]
```

```
##      upper      mean      lower
## 0.4204682 0.4030934 0.3857186
##
## [[5]]
##      upper      mean      lower
## 0.5843593 0.5543236 0.5242879
##
## [[6]]
##      upper      mean      lower
## 0.7213141 0.6941354 0.6669567
##
## [[7]]
##      upper      mean      lower
## 0.8669103 0.8376090 0.8083077
##
## [[8]]
##      upper      mean      lower
## 1.0068237 0.9829162 0.9590087
##
## [[9]]
##      upper      mean      lower
## 1.173365 1.142200 1.111035
##
## [[10]]
##      upper      mean      lower
## 1.324710 1.293345 1.261980
##
## [[11]]
##      upper      mean      lower
## 1.487620 1.438128 1.388636
```

```
df %>% group_by(Size) %>% group_map(~ MY_CI(.x$TimeSeq))
```

```
## [[1]]
##      up      mean      down
## 3.788854e-07 2.000000e-07 2.111456e-08
##
## [[2]]
##      up      mean      down
## 0.1261190 0.1253096 0.1245002
##
## [[3]]
##      up      mean      down
## 0.2664681 0.2637116 0.2609551
##
## [[4]]
##      up      mean      down
## 0.4086907 0.4030934 0.3974961
##
## [[5]]
##      up      mean      down
## 0.5639995 0.5543236 0.5446477
##
## [[6]]
##      up      mean      down
```

```
## 0.7028910 0.6941354 0.6853798
##
## [[7]]
##      up      mean      down
## 0.8470484 0.8376090 0.8281696
##
## [[8]]
##      up      mean      down
## 0.9906179 0.9829162 0.9752145
##
## [[9]]
##      up      mean      down
## 1.152239 1.142200 1.132160
##
## [[10]]
##      up      mean      down
## 1.303449 1.293345 1.283241
##
## [[11]]
##      up      mean      down
## 1.454072 1.438128 1.422184
```

```
df %>% group_by(Size) %>% group_map(~ CI(x=.x$TimePar,ci=.95))
```

```
## [[1]]
##      upper      mean      lower
## 0.0005607236 0.0004596000 0.0003584764
##
## [[2]]
##      upper      mean      lower
## 0.1958966 0.1876130 0.1793294
##
## [[3]]
##      upper      mean      lower
## 0.3437975 0.3373904 0.3309833
##
## [[4]]
##      upper      mean      lower
## 0.519148 0.496546 0.473944
##
## [[5]]
##      upper      mean      lower
## 0.671344 0.655914 0.640484
##
## [[6]]
##      upper      mean      lower
## 0.8253192 0.8020116 0.7787040
##
## [[7]]
##      upper      mean      lower
## 0.9678537 0.9399788 0.9121039
##
## [[8]]
##      upper      mean      lower
## 1.165047 1.105280 1.045514
```



```
##
## [[9]]
##      upper      mean      lower
## 1.360164 1.295007 1.229850
##
## [[10]]
##      upper      mean      lower
## 1.475232 1.428711 1.382191
##
## [[11]]
##      upper      mean      lower
## 1.840906 1.653661 1.466416
df %>% group_by(Size) %>% group_map(~ MY_CI(.x$TimePar))
```

```
## [[1]]
##      up      mean      down
## 0.0004921768 0.0004596000 0.0004270232
##
## [[2]]
##      up      mean      down
## 0.1902816 0.1876130 0.1849444
##
## [[3]]
##      up      mean      down
## 0.3394544 0.3373904 0.3353264
##
## [[4]]
##      up      mean      down
## 0.5038272 0.4965460 0.4892648
##
## [[5]]
##      up      mean      down
## 0.6608848 0.6559140 0.6509432
##
## [[6]]
##      up      mean      down
## 0.8095201 0.8020116 0.7945031
##
## [[7]]
##      up      mean      down
## 0.9489586 0.9399788 0.9309990
##
## [[8]]
##      up      mean      down
## 1.124534 1.105280 1.086027
##
## [[9]]
##      up      mean      down
## 1.315997 1.295007 1.274017
##
## [[10]]
##      up      mean      down
## 1.443698 1.428711 1.413725
##
```

```
## [[1]]
##      up      mean      down
## 1.713982 1.653661 1.593340
df %>% group_by(Size) %>% group_map(~ CI(x=.x$TimeBuiltIn,ci=.95))
```

```
## [[1]]
##      upper      mean      lower
## 9.128245e-06 4.600000e-06 7.175507e-08
##
## [[2]]
##      upper      mean      lower
## 0.2015295 0.1953190 0.1891085
##
## [[3]]
##      upper      mean      lower
## 0.4205594 0.4069590 0.3933586
##
## [[4]]
##      upper      mean      lower
## 0.6702605 0.6335274 0.5967943
##
## [[5]]
##      upper      mean      lower
## 0.9002700 0.8593898 0.8185096
##
## [[6]]
##      upper      mean      lower
## 1.119990 1.082252 1.044513
##
## [[7]]
##      upper      mean      lower
## 1.358966 1.324733 1.290501
##
## [[8]]
##      upper      mean      lower
## 1.651571 1.575591 1.499611
##
## [[9]]
##      upper      mean      lower
## 2.002522 1.841275 1.680029
##
## [[10]]
##      upper      mean      lower
## 2.188603 2.064110 1.939617
##
## [[11]]
##      upper      mean      lower
## 2.508309 2.345241 2.182173
```

```
df %>% group_by(Size) %>% group_map(~ MY_CI(.x$TimeBuiltIn))
```

```
## [[1]]
##      up      mean      down
## 6.058767e-06 4.600000e-06 3.141233e-06
```


```
##
## [[2]]
##      up      mean      down
## 0.1973197 0.1953190 0.1933183
##
## [[3]]
##      up      mean      down
## 0.4113404 0.4069590 0.4025776
##
## [[4]]
##      up      mean      down
## 0.6453609 0.6335274 0.6216939
##
## [[5]]
##      up      mean      down
## 0.8725593 0.8593898 0.8462203
##
## [[6]]
##      up      mean      down
## 1.094409 1.082252 1.070094
##
## [[7]]
##      up      mean      down
## 1.335761 1.324733 1.313705
##
## [[8]]
##      up      mean      down
## 1.600068 1.575591 1.551114
##
## [[9]]
##      up      mean      down
## 1.893221 1.841275 1.789330
##
## [[10]]
##      up      mean      down
## 2.104216 2.064110 2.024005
##
## [[11]]
##      up      mean      down
## 2.397773 2.345241 2.292709
```

As you can see, I was not able to reproduce the behavior of the `CI` function from the `Rmisc` package. I found the source code for the `CI` function [here](#), and I don't understand what they are doing with the `qt` function. Maybe it's related to what *Arnaud Legrand* said about the sample variance being unreliable. However I'm pretty close to the same confidence interval nonetheless.

Exploiting the metadata

I only had time to look at the temperature.

```
df %>% ggplot(aes(x=Size, y=TEMPERATURE)) + geom_point()
```



I believe that what we see here is the first runs of the test slowly increasing the temperature, and then it mostly stagnates around 80-90°.

In general, all the metadata should be able to tell us if something goes wrong.

Using the profiler to figure out in what function most of the code spends its time

I added a new entry in the makefile to compile with the right option in order to use the `gprof` profiler

```
cat src/Makefile
```

```
## parallelQuicksort: parallelQuicksort.o
##
##
##
## CFLAGS = -Wall -O3 -pthread -lrt -std=c99
##
## PROFFLAGS= -pg
##
## %: %.o
## $(CC) $(INCLUDES) $(DEFS) $(CFLAGS) $^ $(LIBS) -o $@
##
## %.o: %.c
## $(CC) $(INCLUDES) $(DEFS) $(CFLAGS) -c -o $@ $<
##
## clean:
```

```

## rm -f gmon.out parallelQuicksort profiling *.o *~
##
## profiling:
## $(CC) $(CFLAGS) $(PROFFLAGS) *.c -o $@
## ./profiling > /dev/null
## gprof profiling gmon.out > gprof.txt
make -C src/ clean
make -C src/ profiling

## make : on entre dans le répertoire « /home/benjamin/git/M2R-ParallelQuicksort/src »
## rm -f gmon.out parallelQuicksort profiling *.o *~
## make : on quitte le répertoire « /home/benjamin/git/M2R-ParallelQuicksort/src »
## make : on entre dans le répertoire « /home/benjamin/git/M2R-ParallelQuicksort/src »
## cc -Wall -O3 -pthread -lrt -std=c99 -pg *.c -o profiling
## ./profiling > /dev/null
## gprof profiling gmon.out > gprof.txt
## make : on quitte le répertoire « /home/benjamin/git/M2R-ParallelQuicksort/src »
cat src/gprof.txt

## Flat profile:
##
## Each sample counts as 0.01 seconds.
## % cumulative self self total
## time seconds seconds calls us/call us/call name
## 80.00 0.32 0.32 1224336 0.26 0.26 partition
## 15.00 0.38 0.06 compare_doubles
## 5.00 0.40 0.02 797 25.09 426.32 quicksortHelper
##
## % the percentage of the total running time of the
## time program used by this function.
##
## cumulative a running sum of the number of seconds accounted
## seconds for by this function and those listed above it.
##
## self the number of seconds accounted for by this
## seconds function alone. This is the major sort for this
## listing.
##
## calls the number of times this function was invoked, if
## this function is profiled, else blank.
##
## self the average number of milliseconds spent in this
## ms/call function per call, if this function is profiled,
## else blank.
##
## total the average number of milliseconds spent in this
## ms/call function and its descendents per call, if this
## function is profiled, else blank.
##
## name the name of the function. This is the minor sort
## for this listing. The index shows the location of
## the function in the gprof listing. If the index is
## in parenthesis it shows where it would appear in

```

```
## the gprof listing if it were to be printed.
##
```

```
## Copyright (C) 2012-2021 Free Software Foundation, Inc.
##
```

```
## Copying and distribution of this file, with or without modification,
## are permitted in any medium without royalty provided the copyright
## notice and this notice are preserved.
##
```

```
## Call graph (explanation follows)
```

```
##
##
```

```
## granularity: each sample hit covers 4 byte(s) for 2.50% of 0.40 seconds
```

```
##
```

## index	## % time	## self	## children	## called	## name
##					<spontaneous>
## [1]	85.0	0.00	0.34		parallelQuicksortHelper [1]
##		0.02	0.32	797/797	quicksortHelper [2]
##		0.00	0.00	864/1224336	partition [3]
##					-----
##				131390	quicksortHelper [2]
##		0.02	0.32	797/797	parallelQuicksortHelper [1]
## [2]	84.9	0.02	0.32	797+131390	quicksortHelper [2]
##		0.32	0.00	1223472/1224336	partition [3]
##				131390	quicksortHelper [2]
##					-----
##		0.00	0.00	864/1224336	parallelQuicksortHelper [1]
##		0.32	0.00	1223472/1224336	quicksortHelper [2]
## [3]	80.0	0.32	0.00	1224336	partition [3]
##					-----
##					<spontaneous>
## [4]	15.0	0.06	0.00		compare_doubles [4]
##					-----

```
## This table describes the call tree of the program, and was sorted by
## the total amount of time spent in each function and its children.
```

```
##
```

```
## Each entry in this table consists of several lines. The line with the
## index number at the left hand margin lists the current function.
## The lines above it list the functions that called this function,
## and the lines below it list the functions this one called.
```

```
## This line lists:
```

```
## index A unique number given to each element of the table.
```

```
## Index numbers are sorted numerically.
```

```
## The index number is printed next to every function name so
```

```
## it is easier to look up where the function is in the table.
```

```
##
```

```
## % time This is the percentage of the 'total' time that was spent
## in this function and its children. Note that due to
## different viewpoints, functions excluded by options, etc,
## these numbers will NOT add up to 100%.
```

```
##
```

```
## self This is the total amount of time spent in this function.
```

```

##
##      children      This is the total amount of time propagated into this
##      function by its children.
##
##      called      This is the number of times the function was called.
##      If the function called itself recursively, the number
##      only includes non-recursive calls, and is followed by
##      a '+' and the number of recursive calls.
##
##      name        The name of the current function. The index number is
##      printed after it. If the function is a member of a
##      cycle, the cycle number is printed between the
##      function's name and the index number.
##
##
## For the function's parents, the fields have the following meanings:
##
##      self        This is the amount of time that was propagated directly
##      from the function into this parent.
##
##      children      This is the amount of time that was propagated from
##      the function's children into this parent.
##
##      called      This is the number of times this parent called the
##      function '/' the total number of times the function
##      was called. Recursive calls to the function are not
##      included in the number after the '/'.
##
##      name        This is the name of the parent. The parent's index
##      number is printed after it. If the parent is a
##      member of a cycle, the cycle number is printed between
##      the name and the index number.
##
## If the parents of the function cannot be determined, the word
## '<spontaneous>' is printed in the 'name' field, and all the other
## fields are blank.
##
## For the function's children, the fields have the following meanings:
##
##      self        This is the amount of time that was propagated directly
##      from the child into the function.
##
##      children      This is the amount of time that was propagated from the
##      child's children to the function.
##
##      called      This is the number of times the function called
##      this child '/' the total number of times the child
##      was called. Recursive calls by the child are not
##      listed in the number after the '/'.
##
##      name        This is the name of the child. The child's index
##      number is printed after it. If the child is a
##      member of a cycle, the cycle number is printed
##      between the name and the index number.

```

```

##
## If there are any cycles (circles) in the call graph, there is an
## entry for the cycle-as-a-whole. This entry shows who called the
## cycle (as parents) and the members of the cycle (as children.)
## The '+' recursive calls entry shows the number of function calls that
## were internal to the cycle, and the calls entry for each member shows,
## for that member, how many times it was called from other members of
## the cycle.
##

## Copyright (C) 2012-2021 Free Software Foundation, Inc.
##
## Copying and distribution of this file, with or without modification,
## are permitted in any medium without royalty provided the copyright
## notice and this notice are preserved.
##

## Index by function name
##
##      [4] compare_doubles      [3] partition      [2] quicksortHelper

```

We see that most of the time is spent in the `partition` function.